

SOFTENG 370 Notes

Theodore Oswandi

October 30, 2017

Contents

1	Lecture 1	4
1.1	Generics	4
1.2	Approaches to Understanding	4
1.3	Usable vs Efficient	4
1.4	OS themes	4
1.5	OS design	5
1.5.1	Themes	5
1.5.2	MS-DOS	5
1.5.3	Early Unix	5
1.5.4	THE Multiprogramming System	6
1.5.5	WinNT and Client/Server	6
2	Lecture 2: History of OS	6
2.1	Total Control	6
2.2	Properties of old OS	6
2.3	Progression: Operators & Offlining	7
2.4	Changes in Hardware	7
2.5	Multiprogramming	7
2.6	Batch Systems	8
3	Lecture 3: History Continued	8
3.1	Scheduling	8
3.2	Power to the people	9
3.3	Time Sharing System	9
3.4	1980s computers	9
3.5	1980s Networking	9
3.6	Multiprocessor Systems	9
3.7	Realtime System	10
3.8	Pocket Computer & Smartphones	10
4	Assignment Notes	10
5	Lecture 4: Virtual Machines	11
6	Lecture 6	13
7	Lecture 7	14
7.1	Runnable	14
7.2	Multitasking	14
7.3	Context Switch	14
7.4	Returning to Running	14
7.5	OTHER STATES	15
7.5.1	Waiting	15

8	Lecture 8	15
8.1	Scheduling Processes/Threads	15
8.2	Levels of Scheduling	15
8.3	Scheduling Algorithms	16
8.3.1	FCFS - First Come First Served	16
8.3.2	Round Robin	16
8.3.3	Minimising Average Wait Time	16
8.4	Handling Priorities	17
8.5	Multiple Queues	17
8.6	UNIX processor Scheduling	17
8.7	Old Linux Process Scheduling	17
8.8	Linux Real-time Scheduling	17
9	Lecture 9	18
9.1	Scheduling with Priorities	18
9.2	Priority Allocation	18
9.3	Theory	18
10	Lecture 10	18
10.1	Problem of Concurrency	18
10.1.1	Example of contention	19
10.2	Critical Sections	19
10.3	Software Solutions	19
10.3.1	Peterson's Solution	19
10.3.2	Bakery Algorithm	19
10.3.3	Interrupt Priority Level	20
10.4	Using Hardware - Test and Set	20
11	Lecture 11	20
12	Lecture 12	21
13	Lecture 13	22
14	Lecture 14	24
14.1	Sockets	24
15	Lecture 15	24
16	Lecture 16	24
16.1	Design Decisions	24
16.2	File attributes	24
16.3	File Name Limits	24
16.4	File Type	24
17	Lecture 17	26
17.1	Cycles	26
17.2	deletion	26
17.3	What's in directory entry	26
17.4	Terminology	26
17.5	Finding File Blocks/Metadata	26
17.5.1	Contiguous	26
17.5.2	Linked Allocation	27
17.5.3	MS-DOS & OS2 FAT	27
18	Lecture 20	27
19	Lecture 21	28

20 Lecture 22 - Memory	28
21 Lecture 23	29
21.1 Half speed memory	29
21.2 Average access time	29
21.3 TLB coverage	29
21.4 Page table size	29
21.5 Inverted page tables	29
21.6 Paging and segmentation	29
21.7 Programs larger than memory	29
21.8 Does it all have to be there?	29
21.9 Locality of reference	29
22 Lecture 24	29
22.1 Effective Access Time	29
22.2 How often we want page faults	30
22.3 Reducing page faults	30
22.4 Working sets	30
22.5 Page Fault Frequency	30
22.6 Choosing pages to replace	30
22.7 Selection algorithms	30
22.8 Windows VMM	31
23 Lecture 25	31
24 Lecture 31	32

1 Lecture 1

1.1 Generics

Operating System The software that makes the computer usable. Using modern computers without an OS is "impossible"

Examples: Windows, OSX, Linux, Unix, iOS, Android, etc...

1.2 Approaches to Understanding

Minimalist

- mostly going to be using this one
- OS contains minimum amount of software to function
- archlike

Maximalist

- All software comes with standard OS release.
- Contains many utilities and programs.
- ubuntuish

1.3 Usable vs Efficient

- make sure you make OS suited for needs
- either specialised or more general purpose
- Think of who you expect to use the system
- If creating a realtime system with potentially thousands of operations in a short amount of time, have to consider efficiency
- Same with battery life if you expect the system to be used in a mobile setting.

1.4 OS themes

Manager Model

- OS is collection of managers, ensuring proper use of devices.
- Managers are independent.
- look out for everything associated with computer
- tie in with hardware. Current state of HW lets OS do more/less things

Onion Model

- Onions have layers (Abstractions)
- resources contained in lower layers.
- Lower layers can't access higher level layers but other way around possible
- Very difficult to get these layers 'right'
- can use in terms of security. Very good idea

Resource Allocator Model

- similar to manager model
- emphasis on fairness and providing services

Dustbin Model

- contains middleware that not considered part of OS
- Sees OS as bits no-one wants to do

Getting Work Done Model

- Idea of it is we use computers to do something else.
- Goal for OS is to help be able to get it all done.

1.5 OS design

1.5.1 Themes

All in one

- All OS components freely interact with each other
- MS-DOS and Early Linux

Separate Layers (Onion Model)

- Simplify verification and debugging
- Correct design difficult to get

Modules

- All in one with modules for some features
- Linux and Windows.

Microkernels

- Client/Server model
- make OS as small as possible
- **Exokernel** puts kernel outside. OS's job only need to authenticate people to use hardware.

VMs

- Java is an example of this

1.5.2 MS-DOS

- Written to provide the most functionality in the least amount of space
- not divided into modules
- Something exokernels trying to do. Make application program access hardware directly.

1.5.3 Early Unix

- UNIX OS in 2 parts. **Kernel** and **System Programs**
- Provides:
 - File System
 - CPU scheduling
 - Memory management
 - Other OS functions
- Ken Thompson and Dennis Ritchie
- Make OS as simple as possible.
- Simple 2 letter commands.
- Ideas of pipelining and process communication

1.5.4 THE Multiprogramming System

- THE was the first to use the layered system
- Contains 6 layers:
 - 5 User programs
 - 4 Input/Output buffering
 - 3 Operator-Console device driver
 - 2 Memory Management
 - 1 CPU scheduling
 - 0 Hardware

1.5.5 WinNT and Client/Server

- WinNT still being still run
 - Win10 now has Windows Subsystem for Linux
- NT provide env subsystem to run code written for differnt OS
- NT and successors are hybrid systems. Parts are layered but some merged to improve performance.

2 Lecture 2: History of OS

- Started at mainframes.
 - Early PDAs were similar to mainframes. Had no memory protection.
- Then go to Minicomputers
- And then desktop
- And how handheld computers

Each of these stages go through cycle of:

1. No software
2. Compulers
3. Multiuser
4. Networked
5. Clustered
6. Distributed Systems
7. Multiprocessor & Fault tolerant.

2.1 Total Control

- Computers expensive in 50s. Data and programs were saved on paper tape.
- Programmers knew how the computer worked. They were very knowledgable about computers.
 - Prepared program and data cards
 - do setup
 - control computer
 - debug
- Computers did 10,000s instructions per second, but were idle a lot of the time.

2.2 Properties of old OS

- **IO polling**, since no other programs running in background, therefore just waiting on input and able to just poll.
- No file system
- No memory management or security
- OS defined by decisions made by user.
- Single program at a time

2.3 Progression: Operators & Offlining

Operators

- Goal is to reduce the time CPU was doing nothing.
- Operators now just "use" the computer. No need for programmer.
 - If something crashes, then just start the next program.
 - Batch similar jobs together, maximise usage of computer.

Offlining

- Form of parallelism in early computing.
- With Big Expensive Computer BEC, but they are just waiting for IO a lot of time. Therefore want to make IO as fast as possible.
- Use smaller computers to convert slower paper to faster magnetic tape. Then that magnetic tape is used as IO for the BEC
- This is the same for output. Have another smaller cheaper computer offload the output magnetic tape from BEC to a printer.

Resident Monitor

- Keep some code in memory.
- It did the work that some operators were doing.
 - clearing memory
 - reading start of new program that needs to be loaded.
 - Can also do some of the IO routines.

Control Programs Standardise the language to communicate with the Resident Monitor. Had tags for things such as \$JOB (for signifying jobs), \$FTN (When fortran compiler needed), and \$END (signifying end of program)

Conclusions from this

- Memory management and file system still not present. Therefore still need to reset if anything bad happens.
- Security patchy at best
- Still need IO polling
- Standard IO routines for programmers
- 2 programs in memory, but one executed
- User interface was JCL (Job Control Language)
- Output of program can be input of another.

2.4 Changes in Hardware

- **Disk drives** provide faster IO.
- Processors that you can **interrupt** also means that there is no more reliance on polling.
- IO devices and CPU concurrent execution, and use local buffer.

SPOOLING (Simultaneous Peripheral Operation On-Line) Meaning that when interrupt, contents of cards read to disk. Therefore current program interrupted.

2.5 Multiprogramming

- Putting multiple programs on at once. Need more memory to do this.
- Now also need for scheduler to manage multiple users' program needs.
 - Need to figure out how to manage stuff. Priority of jobs, how much time to allocate for these jobs, etc...

- No memory protection, so programs could overwrite other program's chunk of memory.
 - Java is an example of something that doesn't give you direct access to memory in JVM.
 - Memory Protection better done by hardware than having software impose limits.
- **Requirements** Limited address range and Operating modes.

Memory Protection Modes

1. User/Restricted Mode
 - Execution is done on behalf of the user.
 - User should not have access to privileged instructions
2. Kernel Mode (SU)
 - Execution done on behalf of the operating system
 - Full access to all instructions.

A **mode bit** can be used to signify what mode a certain program is running in. If something in user mode tries to access memory it is not allocated to, it will go to Kernel mode and throw exception before going back to User mode.

Why we need both We need both because:

If modes existed with relevant instructions, but full memory access; there will still be a lack of memory protection, but also no privilege instruction protection. You can put whatever code you want anywhere.

If memory access limited but no modes or privilege access; then the user will be able to modify amount of memory available for programs.

Memory Protection

- Process gets fixed area of memory that it can use
- If tries to access address out of that range then exception will be thrown.
- Base and Limit register set for each process and how much memory it can have.

2.6 Batch Systems

Memory protection and Processor modes allow you to safely put multiple programs in memory.

Features

- Jobs have their own protected memory
- Disks have file systems. Files linked to owners
- Automated Scheduling. Utilise hardware as much as possible, as operators are slow. Also allows fine tuning of how scheduler works.
- Computer consoles

Not much has changed from programmer's point of view.

3 Lecture 3: History Continued

3.1 Scheduling

- Aims to maximise use of computing machinery OS knows
- Need to know details about device and file processes. What how much resources to allocate.
- Also has to take into account timing and output size.

SOMETHING ABOUT UNIVERSITY OF AUCKLAND SYSTEM

3.2 Power to the people

- Due to hardware becoming cheaper, can have general public own personal computers
- Used to use teletypewriters, but used CRT TVs after a certain point. Editing text was difficult.
- At early 1970s, can code in similar style then you do now.

3.3 Time Sharing System

- People don't like waiting.
 - 200ms+ noticable
 - 5000ms+ unacceptable
- Difficult for scheduler to figure out how to allocate resources. People use different computer differently with differing IO demands.
- Users expect command to run as soon as you press Enter.
- Don't want to have everything run at 100%, otherwise it feels too slow.
- Security an issue for all of these people writing on terminal. Have to increase this and have authentication.

Remnants of Batch Programming

- Has way to run process at given time
- Terminal looked like cards until better graphics came

3.4 1980s computers

- Cycle starts again, started with Resident Monitor Systems.
- Simple single layer file systems
- No security, everything stored on disks. Didn't bother as it was aimed at individual users.
- Did spooling later, for printer output.
- Putting more than one program in memory, using similar system to resident monitor.
- Higher definition screens, pixel addressing for graphics.
- Cycle continues, things like time-sharing features and implementation of UNIX.

Xerox created GUI elements for Office use. Then Apple engineers used ideas to create their Mac.

Features

1. Virtual memory
2. Multiprogramming
3. Complex file system
4. Networking
5. Multi-user

3.5 1980s Networking

Security, Transparency and Protocols/Standardisation create new problems.

Network OS: File sharing, communication scheme, running independent to other machines on network.

Distributed OS: Sharing processing power and resources of lots of computers to make it look like only single system.

3.6 Multiprocessor Systems

Heat is an issue, kind of a soft cap on processor frequency. Therefore can add more cores instead of trying to make each core faster.

Tightly Coupled System Processors sharing memory and clock. Communication through this shared memory. Most computers are now this.

Parallel Systems Mean increased throughput and cheaper way to increase performance. With increased reliability and rate of degradation.

Symmetric Multiprocessing: All core running same OS, most modern systems run this way

Asymmetric Multiprocess Different cores allocated to different jobs/section. Used in very large systems.

3.7 Realtime System

Timing constraints very important.

Hard real-time

- must run within time, or failure happens
- Has to be specifically designed to be hard realtime
- Nuclear plants, air traffic control

Soft real-time

- Doesn't matter too much, more lax.
- Most OS handle soft realtime
- Phone system, multimedia

3.8 Pocket Computer & Smartphones

- Started as PDA/Pocket computers.
- Went through cycle again. Started as resident monitors.
 - Due to hardware limitations, so have to start at the basic level again.
- Battery life and power consumption very important factors.

PalmOS Operating system that PalmPDAs ran on.
Small memory with slow processor.
Efficiency very important factor, to just get passable performance.

Android Popular operating system for current smartphones.
Linux based, application programming in Java.
Google trying to build their own kernel to replace Linux (Fuchsia)

iOS Operating systems that mobile Apple products run on.
Based on OSX (Their desktop OS)
virtual memory and paging for code but not data as writing to flash degrades it.

4 Assignment Notes

- Use standard UNIX symbols to control the threads
- `setupstacktransfer()`
 - **siguser1** represents the user's signal. Let you send stuff to yourself similar to interrupt, but done in software and not hardware
 - **sigaction** is a struct that holds information. Kind of like an object is global due to process having to be able to get to it at any time

- Has a separate, special stack for that signal handler to use.
- **man pages** are really important for this assignment.
- If want to get all man pages relevant to signal then use *man -ksignal*
- Threads need their own stack
 - Running independently of each other and calling their own functions so to guarantee proper functioning it is best for them to have their own stack
- **&setuaction** address of instructions for the signal handler
- **thread1()** contains code that will be executed in the thread
- **threadfunct** is array of names of functions that should be called for all threads If add more then you need to add to the array
- In task 2; 3 threads but 2 of them running the same logic from thread2()
- Information about thread structure found in `littleThread.h`
- static variables aren't allocated on the stack. And preserve value throughout multiple function calls
- MISSED UP TO LIKE 35min in
- **sigaltstack** lets you use that special alternate stack for different threads
 - have malloc some memory and will use it
 - When you call `associateStack()` when making new process you make a new alternate stack
- **kill(getpid(), SIGUSR1);**
 - KILL is system call to get signals. Set it up but haven't associated it with anything it yet
 - KILL sending pid of process you want to send it to. Send signal to yourself (try to kill yourself).
- make local copy of thread in function and set it to READY.
- C doesn't have exception handling. Therefore if error happens in a stack then need the ability to jump to part of memory to give error.
 - **setjmp**: Take snapshot of where you are. Registers of processor (PC will contain this). Can also be used to "freeze" state of a given thread if need to be suspended.
 - **longjmp**: Jump back to state where `setjmp` called. Can be used to "unfreeze" an already suspended thread to resume it. [Line34 in `OSA.c`]
 - Copy stack information/register information and when jump back then recopy it back to "jump back to where you were"
 - variable states preserved if stored on the stack
 - if `setjmp` return 0 then returned directly, or nonzero if from `longjmp`. Will be used later for forking to create new processes, to check if from parent or child
- **Switches** Pass it your current thread and the thread you want to go to.
- Only one thread running at a time, other ones will be READY due to only using a single processor.
- can get this assignment to work without understanding it

5 Lecture 4: Virtual Machines

./ used to signify that it isn't an internal command

MISSED TO VIRTUALISATION

Virtualisation if running on hardware then want to be as close to 90% performance as possible. Preferably 95-98 but not always possible

Design of IBM vm make each user feel like they have own cpu minidisk = lets user feel like they have access to whole drive problem is you don't want actual kernel mode to be accessible to all guests solution is each user has their own virtual kernel mode, but this kernel mode actually runs on the user level. Privileged instructions actually needed to be passed down as not all things kernel does need that mode

Hypervisor Types Allocating resources to VM - like actual CPU cores - or chunks of memory allocated for it Can have "nested" vms

Type 1 Special purpose OS have support for bunch of tools to make using it easier

Type 2 Ones that you install yourself. (virtualbox, parallels) Run applications on host
Problems trap and emulate couldn't be run on x86 up to a point.

Hardware virtualisation x86 Most OS only use level 0 (kernel mode) and level 3 (user mode) Problem with VM in real machine, then you need to keep track about process and registers. Have to keep track of this for all processes. Hardware system lets you change processor for one VM to another

each VM page tables for their own processors used to have nested page table system.
VMs create their own virtual page tables and some will exist in real memory

Solutions Binary translation Look at instructions before execution, problem instructions get translated to be safer to be run in kernel mode

These translations are similar Only translated code is run

OS level virtualisation If lots of machines running same OS, then can use containers that make it seem like they are all separate. Useful for servers Simpler than VMs as they are sharing same copy of OS

More Styles paravirtualisation - XEN modify source code of OS you want to run increase efficiency to allow calls to be made straight to VMM instead of process

Application Virtualisation WINE Want to run something made for an OS on another OS Makes the application think like its running on intended OS

Windows Subsystem for Linux Not really virtual machines If app makes linux kernel call, kernel figures it out and sends it to subsystem Tied into kernel level, applicaiton doesn't really know about it. It just functions as normal and kernel does all of the work.

C and OS implementations

Week 2 friday MISSED 10 MINUTES OF LECTURE

Direct access to memory: address.c Whenever you run the program, the stack address space is different This is for security ASLR address space layout randomisation. Stack, heap and libraries put in different addresses. Helps add level of security

Accessing Registers Can choose to store something in a register Use keyword 'register' prefacing variable type on initialise Can't get address of register, so if set to register OS may put it out of register into memory if you try get address of variable

Volatile Another keyword prefacing variable type Don't do any clever tricks When you don't know if variable value will have changed due to non-local reason due to things like interrupt.

Whenever you use this variable, you have to go back to memory and check its value again as it may have changed.

Memory Management No memory management Static memory allocated at runtime, no malloc. But hard to get rid of them

Dynamic memory Garbage collection doesn't inherently exist. As it is unpredictable

Allocating stack space can be done by calling 'free' Free knows how much memory to free up since malloc uses a little bit more space just above for length of bit of stuff stored

Inline assembly Example code is 32bit OS dependant Can put assembly language directly in C code

Running commands from C program system() lets you put string of command you want to use

Alternatives (languages for OS) C++ similar to C but with object stuff too Windows has C kernel with some C++ and C# Objective C MacOS written with ObjC, but trying to move to Swift Java Can't exclusively use java, need stuff with other stuff as well Assembly old school if you need even more fine tuning

More assignment stuff Part 1 create a lot of threads, link them together (linked list) circular linked list (doubly) keep going around cycle of threads until all finished executing. Only 2 threads given, but should be able to do with n threads Part 2 Add thread.yield() This will call transition system like in part 1. Stop current thread (not finished) and pausing itself to allow another thread to run. Part 3 Interrupt the thread with external source use set.itimer, send signal to processor to signify event happening. (timer has run out, every 20ms) Tells current thread to pause externally and start next thread.

Processes Instance of program execution Thing OS uses as construct to control work

Two parts Resources/Task/Job files open and using windows on screen restrictions on process Code that's running what process is actually doing these days have threads for multiple streams of instructions

Thread sequence of instructions executing without interruption this does happen, but not from thread's point of view. Thread can't tell if it has been paused or not Can run multiple threads but share resources

Typical uses split work accross processors/cores thread for user response, another for some computation task GUI threads and process threads Server applications, have threads for clients. Server preallocates set of threads for handling requests

Thread implementations user level OS sees one thread per process

advantages work if os doesn't support threads easier to make, no system calls application specific control switching is easier (some have register files for threads)

System level operating system knows about it controlled by system calls System knows about state of thread as well. Therefore will schedule based on their state

advantages Threads treated separately If multiprocessor, then can schedule different threads on differnt processors thread blocking in kernel doesn't stop all thread on same process for example if doing read on file, usually code will wait for result and therefore block can allocate cpu to do something in the meantime in this case

Jacketing Check "will I block" before doing something that may block check to see if data already exist in memory. If already there can just get it without having to block. if have to get it, then let processor do something else while it tries to get data

Best of Both worlds Solaris had both system and user level threads before ver9 Uses one to one mapping of user level to kernel level threads. Mapping of single lightweight process to kernel threads. lightweight processes If something on user level thread makes blocking call, other threads on that lightweight process gets to do its thing system makes its own kernel thread and new lightweight process to allow this to happen Windows 7 threads Since Win7 then have user mode scheduling. This also tries to get the best of both world Linux threads Used to not have threads, everything put on one thread Clone call makes a new process. Shares memory, open files and signal handlers Saw them as processes and not threads, so scheduled them Can't signal whole thread, therefore since cloned you aren't sending it to all of them and only the one you specify Killing threads dangerous, due to them sharing memory, then if killed then blocking may cause memory to be in inconsistent state as lock has not been released yet In POSIX, don't actually kill threads. You tell it to cancel itself instead, telling it to die at some point Threads are written in such a way that before it makes a blocking system call, it does some tidying so cancellations can happen Cloned threads can't block if other clone made blocking system call

Week 3 Wednesday More on threads and processes Part 3 assignment numthreads constant will be correct can initialise arrays with that size if you want

6 Lecture 6

Process Control Blocks Things os should know about process BIG LIST GOES HERE process state turns out to be thread state priority used by scheduler owner - security considerations process generally on one processor (306 core moving cost) process group - proceessors working together memory and resource considerations see if process result can be piped to another process, or that it is waiting for result of this process

UNIX process parts can be scattered as parts somewhere else process structure some of information of process held here

user structure not instant access to this in user space some of information of process held here

In UNIX, text = code

WIndows NT split it up into lots of things in ANOTHER BIG LIST TO COPY YAY MISSED SLIDE 3 TO END OF LECTURE EMPHASIS ON FORKS

7 Lecture 7

7.1 Runnable

- On one core, only one thread/process at the same time. (Exception SMT)
- Other processes/threads may be ready to run, or already running

7.2 Multitasking

Pre-emptive Multitasking

- OS uses some kind of criteria to determine how large of a time slice that task
- The more you call yield and switch processes, the more time is wasted and less actual work is done by CPU

Cognitive multitasking Threads know that may have `thread.yield()` called on it and therefore are coded in a way such that when `yield()` is called, issues are less likely to occur

Advantages

- – Control
- – Predictability

Disadvantages

- Critical Sessions
- Efficiency

Co-operative Multitasking

- Two main ways to approach
 1. Process yields right to run
 2. System stop process when system call made
- Doesn't mean task won't run and complete in one go.
- Old UNIX (before 2.6) didn't allow pre-emptive calls when making system calls
 - pre-emptive multitasking always at user level
 - hasn't always been preemptive at system level
 - Actually used to be cooperative in the past
 - Unix was written simpler in the past, expected it to be simple with blocking calls made.

7.3 Context Switch

- Change from one process running to another on same processor, or to handle an interrupt
- Has to save the process state before this can occur
- Context changes as process executes
- Context contains:
 1. Registers
 2. Memory (dynamic elements like call stack)
 3. Files & Resources
 4. Caches

7.4 Returning to Running

State Transition

- Store process properties so it can begin again where it left off
- Page table to be updated if changing processes
- Environment must be restored
- If changing threads on same process then may can just restore registers
- If system has multiple register sets then could thread change with 1 instruction

7.5 OTHER STATES

7.5.1 Waiting

Waiting To stop unnecessary resource consumption Status changed from running to waiting
Suspended Different form of waiting

Java Always had Threads from the start Threads have generally been user level Although Thread.suspend existed that froze thread on system level Thread.resume() to restore it Issue was some resources are tied to one process, and therefore gets a lock Therefore if frozen then other threads can't access it Threads.stop() kill thread and force it to release locks that it may have But may cause data to be left in inconsistent state

Waiting in UNIX WCHAN can contain numbers, represents address in kernel Uses a queue to create a queue for processing Queue associated with hash value or kernel address
HERE GOES SOME PROCESS OF HOW IT ALL WORKS

Finishing Resources used by process need to be accounted for Shared resources usage lowers due to process finishing Make sure tidying up is done, if not done already Don't rely on this, should do this yourself

"When you log out, you want all your processes to finish too" Create a cascading effect, one process shutting down causes other ones associated with it to shut down too

Reasons to Stop Normal Stop must call exit routine does all required tidyup

Forced Stop Only want some processes to be able to kill specific processes. Parents can kill children Children can "generally" kill parents since same owner

UNIX stopping Has 'zombie states' Process that is finished, until parent checks exit status This is a return state/value of a process Used so next process/processes can find out how child finishes and continue execution based on this result If parent is around and child finishes, child becomes a zombie

If parent never calls wait if parent finishes then zombie is freed

Another FSM

Info from Linux Process Table NI = nice value can be positive a negative used to change priorities the lower the number, the higher the priority negative numbers are super priority Only SU can change nice values Normal users can only change nice values to positive values RSS = resident set size memory allocated to it TT = teletype TIME = how long process has been running for CMD = actual command that was executed

8 Lecture 8

8.1 Scheduling Processes/Threads

- **CPU burst time:** time takes for thread running to have to wait for some reason
- Basically, the majority of threads stop after processing for not very long time
- Therefore if we stop them frequently it doesn't make too much of a difference as they are probably waiting anyway

8.2 Levels of Scheduling

Batch Systems

1. Very long term scheduler
 - outside OS, more admin level
 - STUFF
2. Long term scheduler
 - Have multiple queues
 - STUFF
3. medium term scheduler
 - Still programmer dependent how its done
 - STUFF

4. short term scheduler
 - Will mainly look at this one
5. Dispatcher
 - Does the switching from thread/process to another

8.3 Scheduling Algorithms

8.3.1 FCFS - First Come First Served

- No wasting time by determining how to allocate
- Use average waiting time and the CPU burst times for processes
- Produces Gantt chart looking thing
- Weight times are when the processes start

8.3.2 Round Robin

- Pre-emptive version of FCFS
 - Still don't let them run to completion
 - Use of pre-empting them and time slices
- Hard to determine what size time slices to allocate
- Some processes are CPU intensive and require longer time slice
 - But if let these processes do its thing, user may feel slowdown.
 - Interactive processes affected by this
- If short time slice then good in terms of interaction as its jumps around to lots of processes.
 - However, CPU intensive tasks take longer to complete
- Still doesn't have concept of priority
- If task takes shorter time than time slice, instantly schedule another task as to not waste CPU cycles.
- Average wait time reduced due to forced time slices
- Making time slices smaller reduces the average wait time

8.3.3 Minimising Average Wait Time

- Need to know how long CPU bursts are

Shortest Job First

- Gets minimum average waiting time
- But don't always know all CPU burst times
- Therefore use an estimation algorithm. Basing it off previous CPU bursts to estimate how long subsequent bursts will approximately be.

Pre-emptive SJF

- Uses arrival time and burst time
- Short it not because of CLK interrupt, but because another processor came in with a shorter CPU burst time.
- Use remaining CPU burst time remaining if trying to determine if you are going to stop and schedule another process
 - If a process has $CPU_{burst} = 7$
 - Something with $CPU_{burst}=4$ comes at time=2
 - Compare 5 (7-2) with incoming $CPU_{burst}=4$
 - Therefore will stop original processor and run new one since $5 > 4$
- If has 2 options with same weight, then up to programmer to choose. Theoretically similar, but in reality will have some weighting choosing one over the other

8.4 Handling Priorities

Explicit Priorities

- If have very low priority, then there is chance that some priorities will never actually run **Starvation**
- SOMETHING GOES HERE

Variable Priorities

- Processes get higher priorities the longer they've existed (aging)
- Solves the starvation problem

8.5 Multiple Queues

- Multiple queues exist for things that require different time slices and CPU cycles
- Kind of a hierarchy of these processors
- Still assumes single processor

8.6 UNIX processor Scheduling

- Every process has priority associated with it
- Priorities are recalculated every second
- Larger number means worse priority. Lower numbers go first
- Can **Nice** a process, adds priority for process (nicer to everybody else)
 - Ordinary users can only nice their own processes, thereby delaying their processing
- Aging exists, priorities get worse the longer they run
 - Worst level exists, so this doesn't continue forever
 - For every process at worst level, are scheduled in round robin
- The longer a process spends waiting, the lower its priority level becomes and therefore higher chance of being executed

8.7 Old Linux Process Scheduling

- Used two process scheduling algorithms
 1. Time sharing algorithm for most processes
 2. Realtime algorithm for absolute priorities hold over fairness
- Processes have different scheduling classes that determine which algorithm to apply
- Uses **prioritised credit based algorithm** for time sharing
 - Process with most credits go first
 - If process running on clock tick, it loses a credit
 - If process hits 0 credits then another process chosen
 - Therefore the more you wait the more credits you get

8.8 Linux Real-time Scheduling

- Linux uses both **FIFO**(First in first out) and **Round-Robin** scheduling.
- In both situations, processes have priority + scheduling class
- Scheduler does process with most priority
 - If equal priority then choose one that has been waiting the longest
 - FIFO processes run until exit or blocked, no pre-empting
- In Round-robin, processes pre-empted after a while and moved to end of queue.
 - Allows

New Linux Processing

9 Lecture 9

MISSED SLIDE 1 & 2

Periodic process

- common that period and deadline are the same
- Deadlines and period may change depending on the workload of the system

Sporadic Processes aperiodic process things can happen at the same if ∞ events can occur at the same time then need to figure out how to allocate it

Cycling Executives Handle periodic processes Prescheduled - know information before power machine, so can schedule Can't pre-empt because schedule already generated Hard to maintain

CE Schedule MAJOR SCHEDULE MINOR CYCLE

9.1 Scheduling with Priorities

Lets you do important tasks first CATCH UP

9.2 Priority Allocation

Fixed

- **Rate monotomics RM**, shorter period means higher priority
- Least compute time LCT, similar to Shortest Job First

Dynamic

- CPU burst times used/useful
- **Shortest completion time**
 - Simnilar to SJF
 - Uses pre-emption, but requires good information about execution time requirement.
 - Schedule, and compare the time required to finish computation of process at every cycle
- **Earliest Deadline**, process with closest deadline goes first
 - Does this every cycle, and compares all processes that want to be sceduled and their respective deadlines
 - Add don't cares/idle times for when process is complete before deadline. Counts as ∞ , allowing another process to run
- **Least Slack Time** (deadline - compute time) gets highest priority
 - If no slack time left, then must schedule now.
 - Slack time doesn't change if it gets process. Due to fact that its deadline gets closer, but its computation has gone for another cycle, cancelling each other out.

9.3 Theory

- Static priorities, RM is optimal policy
- Dynamic priorities, EDF(Earliest Deadline) and LST(Least Slack) are optimal
- Only really works for single processors.
 - Required more sophisticated processes, to allocate multiple processors

10 Lecture 10

10.1 Prolbem of Concurrency

- Sharing resources is a problem
- Multiple threads/processes trying to access it at the same time

- Some resources can only be safely accessed by a single thread at a time
 - Reading from resources that are being written to.
 - Writing to a file simultaneously
- **Race Condition:** Where order of thread execution produces different results

10.1.1 Example of contention

- Don't have control over thread execution once threads have started executing
- have to `-lpthread` on linux when compiling C programs to use thread library
- `counter++` seems innocent, but actually has a window for error.
 - Due to concurrency, there is contention if multiple threads are the one that calls the function that causes a reaturn
 - Did not count to 10 a lot of the time due to contention happening a lot of times

10.2 Critical Sections

- **Mutual Exclusion:** Area of code that it expects only 1 thread active at given time
- Need to lock thread when critical session going on
- Need to have a way to make sure threads aren't waiting forever (Starvation)
- Starvation can be caused by deadlocks of indefinite postponement

10.3 Software Solutions

- Both can get lock at same time if multithreaded
- Thing dying prematurely
 - OS needs to keep track of when things are alive/dead
 - Always an issue and should have something in place to deal with it, if it were to happen
- Polling puts unnecessary strain on system
- **Spin lock** - something waiting just doing nothing.
- **Busy wait** - Waiting when processor doing something, which is something you don't want to happen
- Suspended thread after seeing that thread is unlocked may lock it at another time, in which another process has already got a lock
- Main problem is multiple threads that sees that `locked = false`, and trying to set it to `true`. Due to small vulnerable time gap

10.3.1 Peterson's Solution

- `flag = [false, false]` is a shared variable
- Getting lock
 - Set your own `flag[self]` to true, and set turn to another thread, and wait until they either they let you or finish doing process on object
 - Setting `flag = true` means that you want to access it.
- Wait when its other thread's turn and they wanna use the file.
- Not feasible in real life due to instruction re-ordering
- Instruction Re-ordering
 - Both compilers and processors do some optimisation at runtime/compile time.
 - Users don't have control over this

10.3.2 Bakery Algorithm

- Each thread given number indicating when it can request a lock.
- Numbers aren't unique so need to use another form of identifier to distinguish different processors with the same allocated number

10.3.3 Interrupt Priority Level

- Can be done by increasing interrupt priority level so other processes can't pre-empt
- Before you do a check on a lock, turn interrupt off so you can't be pre-empted.
- Only turn off interrupt for certain sections of code so that pre-empting will not be possible for that part only and not block other processes.
- Disadvantages:
 - Doesn't work efficiently with multiple processors present
 - A message must be sent to all other processors to let them know of the level change. May cause processors to wait
 - Not all processors at priority level need to be stopped so waste of resources

10.4 Using Hardware - Test and Set

Atomic, create instruction that cannot be divided and must run to completion So the locks that it obtains are guaranteed to not be interrupted. Testing the value, and setting it in one indivisible instruction.

```
while test_and_set(locked)
```

- This removes the gap that is prone to be pre-empted.

This doesn't solve the issue of a busy wait, nor does it make it fair.

Getting out of spin DIDN'T LISTEN YAY

Priority Inversion If low priority process has lock, as long as it can complete and pass the resource on Just need to make sure that the lower priority process can do its task. But sometimes can't due to interrupts and waiting and stuff Don't want lower priority to be stuck in ready state, while it still hasn't been scheduled.

Possible solution temporarily increase priority of something with low priority to higher priority only while it uses the resource. MORE THINGS

Placing in a queue Once new process realises that resource is in use, it suspends, and reschedules itself.

Putting process to sleep then the lock may be set to false so process waiting forever

Put another lock around the code to make sure that locking and unlocking can be done without interrupts

Busy waits solves this issue. OS puts busy waits on lots of little bits of code, but problem with busy waits is it may be waiting for a while.

Semaphores Solution to concurrency problem Basically counter with atomic operations 2 functions, V(S), and P(S). Original value, set to let n number of things NOT SURE AE

P() is kind of a lock If process sees S 0 or less then will wait V() is called when finish May set S ≥ 0 which may let another process use the resource

Implementing Semaphores

UP TO PRODUCER CONSUMER PROBLEM (14) Want consumer to block so that the buffer will have a value stored in it

11 Lecture 11

Reader/Writer problem Readers aren't a problem, as they don't change values As you add writers, then it makes it potentially inconsistent

Only 1 writer at a time in critical section, once it is gone then you can let readers go do their own thing. If we have multiple readers, we want them all in there at once with no writers

Writer preferred writer before reader prioritises reading of updated values read most recent data

Priority problem again, if writers keep coming then readers will wait forever

Reader Preferred Opposite problem to other one, writers may wait forever

both of these may end up indefinite postponement

No preference Use queue, neither writer or reader has no priority

Getting program correct (2) had exclusive access semaphore, set to 1 so that only one thing can access it.

If producer goes twice then overwrites the value in the buffer of original, so no 1:1 mapping of producers to consumers

If consumer gets there before producer then the **number_deposited** is set to 0, producer sees this too and then waits forever and enters a deadlock

Bad Programmers

Have to make sure that if you have a lock in a section of your code, to put something in to unlock it too.

Can use things in OS to help programmers

Can possibly just call unlock if you see another process has a lock on an object and then be allowed to access it. Throws away any safety created if programmer has malicious intents

Monitors Object that allows at most one thread executing inside of it. Similar to old school kernel allowing one process to run in it

As long as you're running in the monitor then you don't have to worry about concurrency issues as the monitor only allows one thing to run at a time.

Slower due to monitor only doing one thing at a time, and therefore potentially getting a big backlog of tasks to do

Condition variables Queue you put thread on when it gets to front Uses the **wait** and **signal** again to only wake threads when needed

If you call wait, you always go to sleep. If consumer sees nothing in buffer then will call wait and sleep itself and wait for producer to be done

If you call signal and thing has changed keep going Otherwise then you do nothing. Used for situation where producers go after another to make sure buffer values are not overwritten.

Which thread runs? If you allow thread to call signal, to wake up another thread then be careful to make sure that you don't interfere with its operation.

Java monitors Every object has a lock. This is to allow **synchronized** methods to work If you call a **synchronized** method then you're asking monitor associated with object if you can go in and do something. Has inherent **wait()** and **signal()** for each object.

What happens when there is a recursive call to **lock()** Throw error if encounters lock by **self** If continue then need to count how many times locked and unlocked. Make sure that you unlock it as many times as you locked it.

pthread_mutex_t Can specify if you want a recursive lock or traditional lock. Traditional lock will not allow same process to lock something twice.

They are different One condition variable for **signal** & **wait** Not real monitors

12 Lecture 12

Dining Philosophers DESCRIPTION

First solution Does **wait()** on both left and right Eat Put them back

They all grab fork on the right and then can't get one on the left. Then die. REST OF THIS

Second Solution Basically the people either have both forks, or no forks Uses simultaneous wait and signal for getting and returning forks No deadlock but doesn't solve starvation problem

REST OF THIS

Simultaneous wait Uses a **try-lock** Pick up one and see if can grab other too, if can't then put one you have picked up back.

Problem with this is it is really slow. The constant contention makes it run a lot slower than first solution, even though that one ran into deadlock It is possible that one single person never eats. A given person can only eat if both its neighbours aren't eating.

Just to be safe 1 THIS

2 THIS

Equivalence If have semaphore then can implement monitor, and vice versa

Implementing semaphore with monitor is easy Other way around is harder due to lack of conditional. Associate semaphore with condition variables.

Lock free algorithms Locks aren't the only way to protect resource/datastructure Libraries exist for this for most languages

Lock free modification uses `cas()` Compare and swap Basically check if value checked has changed. If not then replace with new value. Keep checking until the value

Deadlock When 2 processes want something that the other has and vice versa. (Classic deadlock) Anything waiting for a resource that the 2 processes fighting for also indirectly joins the deadlock.

Conditions

1. Circular linked list
2. Resources can't be shared
3. Only owner can release resource
4. Process holding resource while requesting another. **Often forgotten**

Detection Use graph and find cycles Allocation are node, request are edge

Results Deadlock has to be resolved somehow Killing process. Bad since you don't know current state of process. Can use priority or age to select process to delete. May enter back into deadlock instantly. Can remove all processes. Overkill but solves problem. Can rollback or restart. But make sure that same process not deleted/rolledback constantly

Killing stuff not generally good idea due to potentially inconsistent state.

Prevention Ordering of resources Make circular list of processes, so you can only get resources in a predetermined order. Prevents the basic deadlock situation.

If A- \rightarrow B- \rightarrow C. If have A can get B If have B, have to release, get A then get B

Resources not sharable.

Only owner can release resource

Process can hold resource while requesting another One resource at a time Return before requesting Allocate all resources at same time.

Avoidance More of a runtime thing. Very conservative strategy Check requests' safety, before allowing it to go ahead. Might end up have resource free while something waits for it due it it possibly making deadlock later.

Occur when both processes have resources and both require one more from depleted resource pool.

Banker Algorithm If given request, then assume permission granted. Go through processes, and see what processes it wants not and in the future, compare with the list of available resources. Then actually grant it access if it meets this. Otherwise do nothing.

Check if the processes can finish with given resources left if trying to allocate to another. Need to ensure that all processes can definitely finish. "Deadlock may occur" not allowed. Only "deadlock can't occur" allowed

13 Lecture 13

Distributed Deadlock Instead of worrying about order that you get resources, instead order processes that are executed by their priorities. Lower levels are rolled back and higher ones go earlier

Detection A - \rightarrow B means that A is waiting on something B holds Sometimes that you can only see deadlock when you combine multiple graphs. Can have something ask all sites for their waitfor graphs (to try find global deadlocks)

Centralised deadlock detection Timing is an issue and may lead to false positives. Graph only approximation of real resource allocation. Can use timestamps to avoid false deadlock

Distributed approach Extra node (P_{ex}) in each local waitfor graph Local processes waiting on external stuff goes to this P_{ex} If cycle with P_{ex} in it exist, **could** have a deadlock. Each local site can check this. Then goes and ask other sites Deadlock handled if found Otherwise continue as normal

Time stamp prevention wait-die If resource held by older process, younger process can't wait, and dies. Since it hasn't done as much work, and tries again.

If resource held by younger process, older process allowed to wait.

Keep age for when you first try get a resource

wound-wait Younger process allowed to wait for older process. Other away around. Older process doesn't wait for younger process. Kills it and takes resource.

Messages Can be used to control concurrency Sending information to another process can be done by

1. Shared resource
2. Message passing
 - Address message
 - Transport message
 - Notify receipt ANOTHER ONE
 - `send(destination, message)`
 - `receive(source, message)`
 - `write(message)`
 - `read(message)`

Design decisions Either have sender block or not If writing to a file, want to send information as soon as possible.

Receiver should be blocking, you don't want any interruptions. Receiver can choose a bunch of message types to stick around for and receive if any come.

If sender doesn't block then need somewhere to store the messages until the receiver retrieves the message.

Storing messages

1. Have sender send message straight to other process
2. Or have the page saved, and pointer sent so receiver knows where it is
 - Need to not overwrite information until it has been read
 - Should also make this address read only for recipient so it can't be tampered with

COMMUNICATION PROCESS TO PROCESS (12)

Indirect Communication Mailbox and ports Mailbox ownership

1. Owned by system
 - Persist without process (if finish)
2. Owned by process
 - Creator pass ability to receive
 - mailbox gone if process finish

UNIX pipes General Something that contains data Buffering mechanism Pipe fills designated array with file descriptors (address on open file table) which basically makes it look like it is variable that contains the file/s Can use typical `read` and `write` calls `ls -al | grep Doc | wc -l` Vertical bar is pipe

1. Create new process using fork Call `exec` on `ls` Forks to `grep`
 - `ls` end is writing, `grep` end is reading

UNIX does this piping by redirecting `stdin` and `stdout`

Full Pipes If full (due to specified pipe size), and tries to send more information through the pipe. Will block when buffer fills up to not lose any information

Broken Pipe If trying to read from pipe with no sender, get EOF. Otherwise it may wait forever. If reader has died, sender has no point to keep send information; send signal to the process letting it know that nobody is listening to it.

Pipes aren't simply single sender and single reader. Chunk sizes at least 512bytes. Then if message is equal or smaller than it, it will keep it as a whole.

Mach ports Underneath OSX Ports that write/read to/from. Only one reader per port Process allocated a certain number of ports for communication

UNIX process communication SIGUSR1 user level signal you were sending

14 Lecture 14

14.1 Sockets

15 Lecture 15

16 Lecture 16

16.1 Design Decisions

need to store variable files and different types block is smallest addressable region see it as an array of blocks sector is smallest area the device could communicate with

16.2 File attributes

Information about files things

- Name
 - may have more than 1 name for a file
 - length of name
 - limitations to name (can't use characters)
- Location
- Size
- Owner
 - Protection, and grant access rights
- Access Information
- Dates/Times (Creation, modification)
- File Types

16.3 File Name Limits

- NTFS
 - extend paths up to 32k length
 - ordinary path up to 256chars
 - Win10 can ignore limit
- Linux
 - Path limit 4k bytes
 - commonly 255 bytes per path component
- APFS
 - 255 characters per path component
 - Don't know maximum length
 - iOS 10.3 converted this on devices on the fly.

16.4 File Type

System has to know about it to perform operations

;;EG SECTION TO COPY;;

All OS know about executable binary files OS specific structure about how to load information from file/s

Dealing with File Types

- Easy way to do it - just add file type to the end of the name.
- Extensions are mapped to registry so it knows how to treat it
- Doesn't stop people renaming extension
- UNIX uses numbers in front of file data (first few characters/bytes have this information)
 - `ftype M4A` means that it is .m4a type
 - `PDF1.3` - PDF format

Old Macintosh Solution

Used to be called MFS (Mac File System) Wanted to have icons for files Stored too much information. stored what program created the file so it can be opened again later.

2 parts to a file resource fork `;;LIST OF THINGS;;`

data fork `;;LIST OF THINGS;;`

Done before computers were networked. Things broke due to OS incompatibilities. This is incorporated into NTFS, but extended to have as many parts as you want to a file as programmer wanted.

in OSX, **bundles** seem like a single file, but is actually a directory.

NTFS Highly structured way to look at a file. File seen as a set of file attributes. Only one of the attributes contain the file data.

File system commonly have a Btree to make finding files more efficient.

Alternate Data Streams Hidden information about files Need to use other programs to view it

Starts with colon. Associates this alternate data stream with current directory.

Master file table, each file has approximately 1kb allocated to it.

Representing Files on Disk Have to store files in constant sized blocks on disk

`;;LAST 2 BULLETS;;`

Structure metadata = `;;DEFINITION;;` single level = `;;DEFINITION AND USE;;`

Number of files grow as you put files in, B-tree used sometimes.

Two level users are top level, can only have flat directory not used by many OS

Normal tree structure files organised by using collections of more files

Should you be able to write directory that you own Don't want to lose some of the metadata that contains what is stored inside directory can also lose list of blocks allocated, then you get full access to disk if you modify this

Sharing Files and Directories `;;LOTTA TEXT;;`

Hard Links `;;MORE TEXT;;` With hard links, both files are treated equal.

`ln old_filename new_filename` Create directory entry with pointer to file inode (hold real information about a file)

`mklink /H newfilename oldfilename` How to do it in Windows. /H to denote that it is to be a hard link. Make entry of master file table. But also puts a bunch of information in so you don't have to go to the master file table every single time.

Soft/Symbolic Links The contents of a symlink is the identifier to the real file. OS knows to treat it as link due to the `l` in `ls -l` Length of the file is length of filename that it points to Only holds real name of the file. If you rename file, symlink breaks

Windows shortcuts `;;THING;;` Windows should be able to continue tracking even if you do things to the original file. Such as moving file to different volume, move to network. However this uses lots more resources and does more work to allow this.

Mac Aliases Can make alias from GUI. Contains real identifier of the file.

However, they moved to UNIXlink at some point and incorporated symlinks. But symlinks work differently and are 'dumber'

Now MacOS stores both ID and filename. Checks the file name first, then ID if file doesn't exist.

17 Lecture 17

17.1 Cycles

Generally don't want cycles in the directory graphs Have infinitely many names for all files in the cycle. Naive search would fail, since you may get stuck in the loop and never reach file you're searching for.

One solution would be to just not allow hardlinks, but allow symlinks. Can still get infinite loop, but OS thinks its in loop if it reaches a certain number of symlinks.

17.2 deletion

Have count of how many links to files If delete, then reduce the number so other names can still point to data If real file deleted then could have dangling pointers

17.3 What's in directory entry

UNIX

- filename
- file attributes (How much more?)
- pointer to file attributes

UNIX stores this attribute information in **inode** that also contains number of hard links to file. UNIX directory is a table of names, and their inode numbers

inode table replicated so not a single point of failure.

NTFS Master file table contains folder information.

;;YEP NEED TO TYPE A BIT;;

Not all information you get may be correct

17.4 Terminology

17.5 Finding File Blocks/Metadata

data on disk that isn't contents of file

17.5.1 Contiguous

store start block and length of each file. Efficient for hard disk as head doesn't have to move very much. If have both of these values then can deal with any part of the file.

Can't do simple things like make file larger. Need to find hole large enough in disk to fit new file size.

1. **First fit**
Find first hole that will fit
2. **Next fit**
Continue searching where you last finished searching
3. **Best fit**
Fit as best you can, hole closest in size. Disk ends up with small holes left that you can't fit more files in
4. **Worst fit**
Find hole with largest space, lets you fit more files in there later.
5. **Buddy algorithm**

Figuring out how much space to allocate is also difficult thing to do. If give too little space then increase in file size annoying Otherwise, may be wasting space.

17.5.2 Linked Allocation

Little bit of block used to point to next block that is part of a file. Linked list implementation.

Simple. No more external fragmentation, all blocks can be used.

Direct or random access is difficult. Need to read block to know where next block is. If one block fails, can lose rest of file.

Can cache blocks associated with a given file.

17.5.3 MS-DOS & OS2 FAT

File allocation table. Part of disk holds this table, is array of numbers that represent blocks.

If can load whole allocation table to memory, then can do direct access.

Works best if number of things allocated isn't very large as it is preferable to load whole table into memory.

;;MISSED A LECTURE;;

slide 12;;Opening a file in UNIX;;

fd example doesn't use root directory changing where stuff redirects to/from by closing an opening file.

;;FILE VERSIONING;;

;;19 slide 15;; Better Methods ;; Transparency Things are hidden. Not seeing the details/implementation User shouldn't be able to see differences/complications Location Transparency No connection between name of file and where it is stored

Migration Transparency Moving file without noticing it moving Location Independence - name doesn't change when you move it

Collections of Files Don't keep all information for all files ;;

Using remote files Need to transfer file data over network to client May cache it to client memory or local disk. Can be expensive Also need to let all users know if cached, that somebody has modified the file.

Caching Blocks of file cached locally Accessing file done to the cache Modify file: Write through - every write requires user to send block back to server Delayed Write - Update sent to server delayed, or when file closed.

;;PROS CONS;;

Consistency Semantics way changes in data is distributed

UNIX - changes immediately visible SESSION - process gets copy of file and changes not visible until file is closed.

18 Lecture 20

Stateful Server knows who has file and what they can access. Keeps track of file pointer too. Client gets identifier to use to access the file Server can intelligently cache files that it knows client it already trying to access State information lost if server shuts down/crashes. If client dies,

Stateless Server doesn't hold anything about state of system Need to pass information to and from server constantly. A lot easier to restart the server/system. If client goes down, not a problem from server POV. Since it only just replies to requests. Only if pure stateless. In reality the server will store some kind of information.

NFS Remote service technique, mostly stateless Add some information to the inode. Contains a bit to denote if file is local or not. No need for dedicated servers All machines share files with each other. Every machine can be server Can mount things wherever you want (If you're an admin) Works in heterogeneous environment. Doesn't matter what machines were using, NFS can work. Need to make sure that when you perform operation on server, that adaptations/conversions happen since server and client may not be running in same environment.

Automounter Client mounting/unmounting directories in other systems on demand. ;;MORE;;

NFS protocol make system call virtual file system determine if local or remote ?? request goes to vfs on remote machine does process locally go back to result

;;PROBLEMS WITH NFS;;

AFS mount point at /afs/... tries to cache as much as possible. Uses location database to track where things are.

Shared access to files Access controllers Session semantics - file only updated when file is closed. Callbacks Promise from server that given file is up to date. Once server gets update to file, breaks all callbacks other clients have about the modified file Lost update possible if 2 clients update the same file.

19 Lecture 21

SAN and NAS NAS SAN

20 Lecture 22 - Memory

classic memory hierarchy primary = ram secondary = Disk Tertiary = Archival

Address binding Compile Time Put instruction at fixed address, call using jump instruction

Load Time Object modules uses tables of offsets Mapping done as modules loaded Precompiled libraries

Run time Mapping of final address maintained on the go Used in conjunction with load and compile times

Memory spaces Need to protect OS memory from programs, and enabling programs larger than memory to run

Split memory Can protect with single fence register.

Dividing memory 2 registers to denote how much memory a job should have access to. If try to get something out of that range then should not work since it is violating your allocated space

Two different addresses Make process feel like it starts at address 0. Uses an offset to get real address.(relocation register)

This is a big conceptual change, since it now has 2 addresses to represent the same address.

Everything needs to be in contiguous memory.

Splitting memory to smaller chunks Uses something similar to file allocation table. Stores [Chunk, base, limit]

Two approaches Same sized chunks - pages Variable sized chunks - segments

classic argument about variable and static about flexibility and simplicity

Paged System address translation Local address divided into

1. Page number

Index to which page table contains base address

2. Page offset

Added to base address to get physical address.

Frames and Pages

Tables

Different sized chunks

Segments

Allocation strategies No internal fragmentation, only allocate space we want Can get external fragmentation Strategies

1. First fit

2. Next fit

3. Best fit

4. Worst fit

Defrag memory if need large chunk, faster than doing it to disk

How much space in hole Knuth's 50% rule. If N segments then N/2 holes. If average hole size same as average segment size, need to keep 1/3 of memory free

21 Lecture 23

21.1 Half speed memory

- If paged or segmented memory, then need 2 memory access to get logical memory address.
 - Page/segment table query
 - Data query
- Memory management unit caches recent page table information in faster hardware cache Stored in **Translation look-aside buffer TLB**

21.2 Average access time

Hit ratio - times pages found to not found in associated registers

21.3 TLB coverage

Smaller the cache the better But generally need larger page sizes or overall cache too small. But more IO therefore internal fragmentation Variable page size good, but need good allocation algorithm

21.4 Page table size

- If 32bit address space with 12bit offset, then get about a million entries
- Most process don't use all memory
- Limit page table values to valid ones
- Use page table length register, with valid bit flag
- Only allocate parts of table that we need

21.5 Inverted page tables

Page table now only contains [pid, page number] Frame is the index of the page table. Address based off both pid and page number, hash table can be used for allocation.

21.6 Paging and segmentation

21.7 Programs larger than memory

Swapping - moving pages in and out of disk

21.8 Does it all have to be there?

Try provide more memory than RAM In x86 there may be less logical memory than physically there.

21.9 Locality of reference

22 Lecture 24

22.1 Effective Access Time

Since processors now thinks it has fast memory, but stored on disk. It will appear that the process is a lot slower due to having to access that much slower memory.

Calculation

Accommodate for:

- TLB hit
- TLB miss and page table hit
- TLB miss and page table miss, requiring getting another page.

22.2 How often we want page faults

Don't want it to happen very often. If want to be half as slow (compared to everything in RAM) then need to only get page fault 1/600,000 times

22.3 Reducing page faults

- Different processes access different number of pages in memory.
- Allocate frames equally or proportionally depending on priority.
- Need to have minimum & maximum number of pages.
- Need currently required pages in memory

22.4 Working sets

- Used to keep track of pages needed in real memory to keep process running. Observe process over short window and record the page accesses to determine the working set.
- Finding this time window is hard since you don't want it too short (not enough information) or too long (too many pages)
- Using a **reference bit** on the pages that are accessed allows identification of these used pages.

22.5 Page Fault Frequency

- Used to control number of frames allocated to process
- The frequency of page faults drops a lot when you initially increase number of pages. But law of diminishing returns exists, so set upper and lower bounds and add/remove frames to stay within them.

22.6 Choosing pages to replace

- Need to replace page if no free ones left
- 2 ways to select these pages
 - **Global:** any frame can be chosen. Number of frames for process varies.
 - **Local:** frame must be in process' allocated frame set. There are less frames to choose from.
- Normally global replacement used.

Picking:

- Unmodified pages don't have to be written to disk. **Dirty bit** in page table entry used to determine if changes have been made to frame.
- Pages that aren't going to be used soon in the future should be replaced, but can't see into the future.

22.7 Selection algorithms

1. Random

- Every process treated equally
- Easy to implement
- With enough pages the chance of worst case very low

2. First In First Out (FIFO)

- Queue of pages, remove at head add to tail
- Simple
- Important pages replaced as often as less important ones
- **Belady's anomaly** - increasing number of pages can increase page faults.

3. Least Recently Used (LRU)

- Assumption that if page not used recently then won't be used in near future
- Generally better than FIFO, but not all the time.
- Cannot suffer from Belady's anomaly because recently accessed pages have higher priority and will therefore not be replaced. It does not assume that the first page loaded in is safest to replace.
- More expensive as need to find a way to store last access time for each page.
- **Or** can have list of pages and move to top of list if accessed.
- **Approximations:**
 - **Reference bit** Set when page is used. One per tick, a right shift is done to the value with either a high or low passed in depending on if the page was accessed. These values are compared for each page and one with lowest value is replaced.
 - **Second chance:** Algorithm based off FIFO, but if referenced bit of first in page is set, it unsets it and moves it to back of queue and checks the next one.

4. Least Frequently Used (LFU)

- Uses count of memory accesses per page
- When page needs replacing, does it on page with lowest number
- Pages can stay longer than needed if accessed many times in short period of time.
- Not commonly used

5. Most Frequently Used (MFU)

- Reverse of LFU
- Concept behind this is that pages with few accesses are newly added and may need to be used later
- Not commonly used

22.8 Windows VMM

- **Virtual Memory Manager** runs in background maintaining memory policies.
- Processes use working set min/max
 - Process guaranteed working set min
 - If number of frames below max then try to allocate more
 - If not enough frames then trimmed to min
 - Default size = 30
- Privileged processes lock pages in real memory, for things like real-time processes and device drivers
- **Clustering** - pages surrounding also get brought in. Assume that if you pull in a page then you may need to use ones around it too to reduce page replacement time.
- Windows prefetching
 - pages and files references in 10 seconds of opening application
 - Keep track of this information, loaded up again next time app opened.
 - Defrag every couple of days.

22.9 Thrashing

- When number of pages in all working sets higher than number of available pages.
- Everything will be a page fault, severely affecting work efficiency.

- **Batch system** can thrash if set up to increase number of programs at a time. But low CPU utilisation will lead to thrashing.

22.10 Location of process memory

Addressable memory in UNIX/Windows scattered in many places and page files not loaded yet or paged out.

23 Lecture 25

setuid **chmod u+s** on something. When program is run, runs as it if was executed by user.

If you can create another process when in a setuid program, then will be assumed that uid had executed/created it.

24 Lecture 31